

DowJones_

By Hatem Elgenedy

November 7, 2025

```
[ ]: import numpy as np  
import seaborn as sns  
import pandas as pd  
import matplotlib.pyplot as plt
```

```
[82]: print(sns.get_dataset_names())
```

```
['anagrams', 'anscombe', 'attention', 'brain_networks', 'car_crashes',  
'diamonds', 'dots', 'dowjones', 'exercise', 'flights', 'fmri', 'geyser', 'glue',  
'healthexp', 'iris', 'mpg', 'penguins', 'planets', 'seairce', 'taxis', 'tips',  
'titanic']
```

```
[83]: df = sns.load_dataset("dowjones")  
df.head()
```

```
[83]:      Date  Price  
0  1914-12-01  55.00  
1  1915-01-01  56.55  
2  1915-02-01  56.00  
3  1915-03-01  58.30  
4  1915-04-01  66.45
```

```
[84]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 649 entries, 0 to 648  
Data columns (total 2 columns):  
 #   Column  Non-Null Count  Dtype     
---  --     --          --  
 0   Date    649 non-null    datetime64[ns]  
 1   Price   649 non-null    float64  
dtypes: datetime64[ns](1), float64(1)  
memory usage: 10.3 KB
```

```
[85]: df.shape
```

```
[85]: (649, 2)
```

```
[86]: df.isnull().sum()
```

```
[86]: Date      0  
Price     0  
dtype: int64
```

```
[87]: df.Date
```

```
[87]: 0    1914-12-01  
1    1915-01-01  
2    1915-02-01  
3    1915-03-01  
4    1915-04-01  
...  
644   1968-08-01  
645   1968-09-01  
646   1968-10-01  
647   1968-11-01  
648   1968-12-01  
Name: Date, Length: 649, dtype: datetime64[ns]
```

```
[88]: df.duplicated().sum()
```

```
[88]: 0
```

```
[89]: df.Price
```

```
[89]: 0      55.00  
1      56.55  
2      56.00  
3      58.30  
4      66.45  
...  
644   883.72  
645   922.80  
646   955.47  
647   964.12  
648   965.39  
Name: Price, Length: 649, dtype: float64
```

```
[90]: df.describe()
```

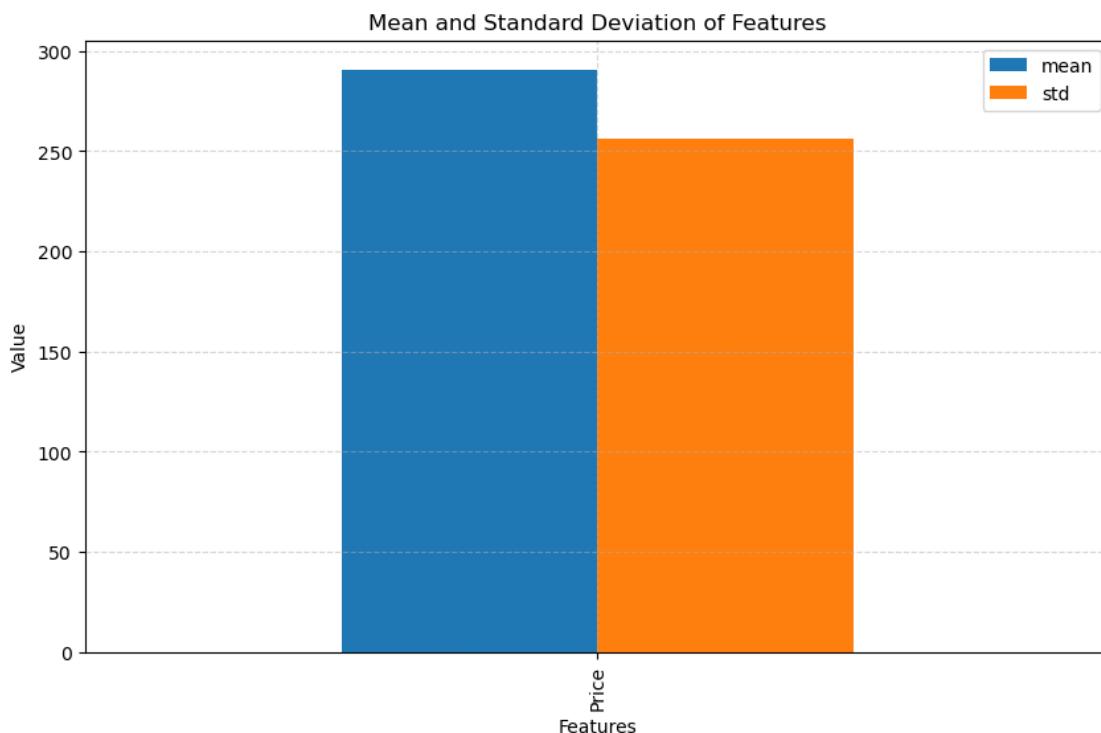
```
[90]:
```

	Date	Price
count	649	649.000000
mean	1941-11-30 17:42:48.258859776	290.807319
min	1914-12-01 00:00:00	46.850000
25%	1928-06-01 00:00:00	106.900000
50%	1941-12-01 00:00:00	172.270000
75%	1955-06-01 00:00:00	436.730000
max	1968-12-01 00:00:00	985.930000

```
std           NaN    256.062906
```

```
[195]: desc = df.describe().T

desc[['mean', 'std']].plot(kind='bar', figsize=(10,6))
plt.title('Mean and Standard Deviation of Features')
plt.xlabel('Features')
plt.ylabel('Value')
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```



```
[91]: len(df.Price)
```

```
[91]: 649
```

```
[92]: len(df.Date)
```

```
[92]: 649
```

```
[93]: df.Date.value_counts()
```

```
[93]: Date
1914-12-01      1
```

```
1952-02-01      1
1950-08-01      1
1950-09-01      1
1950-10-01      1
..
1933-02-01      1
1933-03-01      1
1933-04-01      1
1933-05-01      1
1968-12-01      1
Name: count, Length: 649, dtype: int64
```

```
[94]: df.Price.value_counts()
```

```
Price
98.50      2
97.00      2
120.10     2
99.45      2
75.70      2
..
102.85     1
106.90     1
102.30     1
96.20      1
965.39     1
Name: count, Length: 637, dtype: int64
```

```
[95]: !pip install ydata-profiling
```

```
Requirement already satisfied: ydata-profiling in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (4.17.0)
Requirement already satisfied: scipy<1.16,>=1.4.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (1.13.1)
Requirement already satisfied: pandas!=1.4.0,<3.0,>1.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (2.3.1)
Requirement already satisfied: matplotlib<=3.10,>=3.5 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (3.10.0)
Requirement already satisfied: pydantic>=2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (2.11.7)
Requirement already satisfied: PyYAML<6.1,>=5.0.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (6.0.2)
Requirement already satisfied: jinja2<3.2,>=2.11.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
```

```
profiling) (3.1.6)
Requirement already satisfied: visions<0.8.2,>=0.7.5 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling) (0.8.1)
Requirement already satisfied: numpy<2.2,>=1.16.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (1.26.4)
Requirement already satisfied: minify-html>=0.15.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (0.18.1)
Requirement already satisfied: filetype>=1.0.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (1.2.0)
Requirement already satisfied: phik<0.13,>=0.11.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (0.12.5)
Requirement already satisfied: requests<3,>=2.24.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (2.32.4)
Requirement already satisfied: tqdm<5,>=4.48.2 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (4.67.1)
Requirement already satisfied: seaborn<0.14,>=0.10.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (0.13.2)
Requirement already satisfied: multimethod<2,>=1.4 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (1.12)
Requirement already satisfied: statsmodels<1,>=0.13.2 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (0.14.5)
Requirement already satisfied: typeguard<5,>=3 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (4.4.4)
Requirement already satisfied: imagehash==4.3.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (4.3.1)
Requirement already satisfied: wordcloud>=1.9.3 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (1.9.4)
Requirement already satisfied: dacite>=1.8 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (1.9.2)
Requirement already satisfied: numba<=0.61,>=0.56.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (0.61.0)
Requirement already satisfied: PyWavelets in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from imagehash==4.3.1->ydata-profiling)
(1.9.0)
Requirement already satisfied: pillow in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from imagehash==4.3.1->ydata-profiling)
(11.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from jinja2<3.2,>=2.11.1->ydata-profiling)
```

(3.0.2)

Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (1.3.1)

Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (0.11.0)

Requirement already satisfied: fonttools>=4.22.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (4.55.3)

Requirement already satisfied: kiwisolver>=1.3.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (1.4.8)

Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (24.2)

Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (3.2.0)

Requirement already satisfied: python-dateutil>=2.7 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling) (2.9.0.post0)

Requirement already satisfied: llvmlite<0.45,>=0.44.0dev0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from numba<=0.61,>=0.56.0->ydata-profiling) (0.44.0)

Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling) (2025.2)

Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling) (2025.2)

Requirement already satisfied: joblib>=0.14.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from phik<0.13,>=0.11.1->ydata-profiling) (1.5.1)

Requirement already satisfied: charset_normalizer<4,>=2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (2.5.0)

Requirement already satisfied: certifi>=2017.4.17 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (2025.8.3)

Requirement already satisfied: patsy>=0.5.6 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from statsmodels<1,>=0.13.2->ydata-profiling)

(1.0.2)
Requirement already satisfied: typing_extensions>=4.14.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
typeguard<5,>=3->ydata-profiling) (4.15.0)
Requirement already satisfied: attrs>=19.3.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from
visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
(24.3.0)
Requirement already satisfied: networkx>=2.4 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from
visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
(3.4.2)
Requirement already satisfied: puremagic in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from
visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
(1.30)
Requirement already satisfied: annotated-types>=0.6.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
pydantic>=2->ydata-profiling) (0.6.0)
Requirement already satisfied: pydantic-core==2.33.2 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
pydantic>=2->ydata-profiling) (2.33.2)
Requirement already satisfied: typing-inspection>=0.4.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
pydantic>=2->ydata-profiling) (0.4.0)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from python-
dateutil>=2.7->matplotlib<=3.10,>=3.5->ydata-profiling) (1.17.0)

[96]: `from ydata_profiling import ProfileReport`

[97]: `profile = ProfileReport(df , title = "DowJones Profiling Report" , explorative=True)`

[98]: `profile.to_notebook_iframe()`

Summarize dataset: 0% | 0/5 [00:00<?, ?it/s]
100% | 2/2 [00:00<00:00, 1639.36it/s]
Generate report structure: 0% | 0/1 [00:00<?, ?it/s]
Render HTML: 0% | 0/1 [00:00<?, ?it/s]
<IPython.core.display.HTML object>

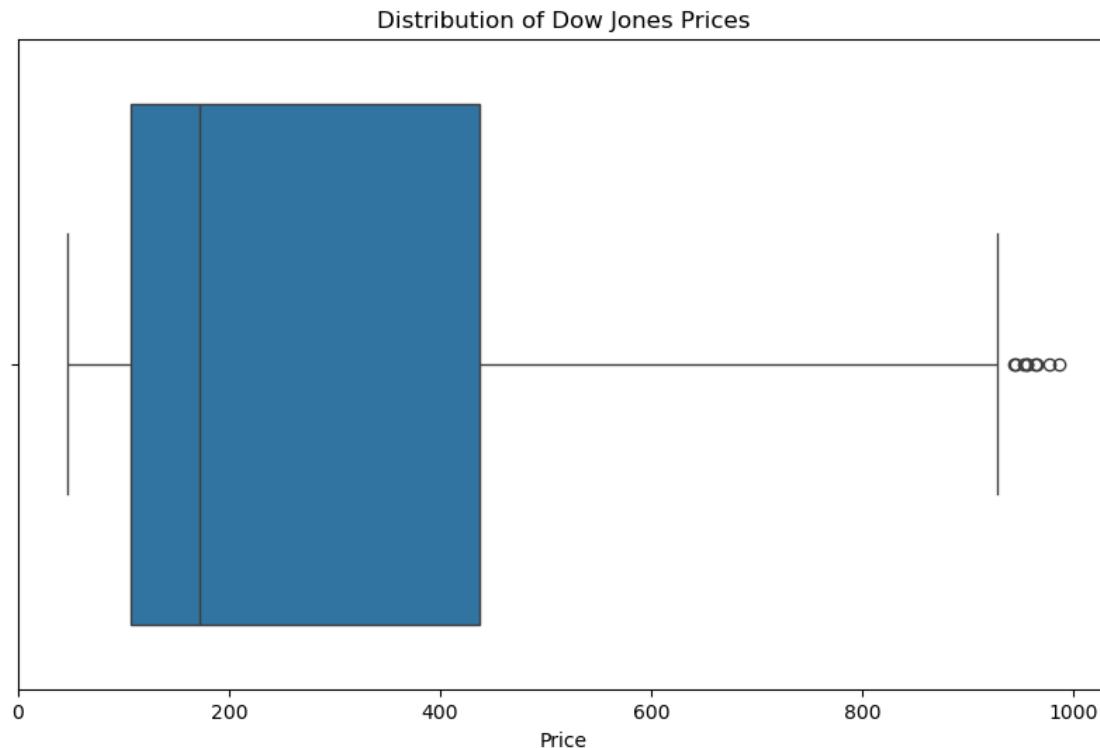
[99]: `profile.to_file("DowJones_report.html")`

Export report to file: 0% | 0/1 [00:00<?, ?it/s]

```
[100]: %matplotlib inline
```

```
fig = plt.figure(figsize=(10,6))
ax = fig.subplots()

sns.boxplot(x=df["Price"], ax=ax)
plt.title("Distribution of Dow Jones Prices")
plt.xlabel("Price")
plt.show()
```

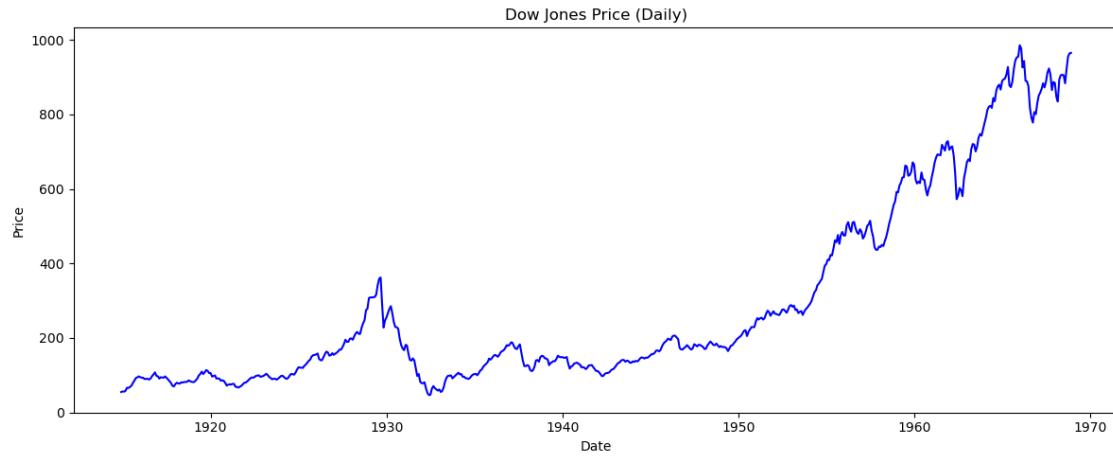


```
[101]: import matplotlib.dates as mdates
```

```
fig, ax = plt.subplots(figsize=(12,5))
ax.plot(df['Date'], df['Price'], color='blue')
ax.set_title('Dow Jones Price (Daily)')
ax.set_xlabel('Date')
ax.set_ylabel('Price')

locator = mdates.AutoDateLocator()
formatter = mdates.ConciseDateFormatter(locator)
ax.xaxis.set_major_locator(locator)
ax.xaxis.set_major_formatter(formatter)
```

```
plt.tight_layout()  
plt.show()
```



```
[102]: df1 = df.copy()
```

```
[103]: df1['Price'].describe()
```

```
[103]: count      649.000000  
mean       290.807319  
std        256.062906  
min        46.850000  
25%       106.900000  
50%       172.270000  
75%       436.730000  
max       985.930000  
Name: Price, dtype: float64
```

```
[104]: q1 = df1['Price'].quantile(0.25)  
q3 = df1['Price'].quantile(0.75)  
print(q1,q3)
```

106.9 436.73

```
[105]: IQR = q3 - q1  
print(IQR)
```

329.8300000000004

```
[106]: upper_limit = q3 + 1.5 * IQR  
lower_limit = q1 - 1.5 * IQR
```

```
[107]: upper_limit
```

```
[107]: 931.4750000000001
```

```
[108]: lower_limit
```

```
[108]: -387.845
```

```
[109]: df1[df1['Price'] > upper_limit]
```

```
[109]:      Date    Price
610  1965-10-01  944.77
611  1965-11-01  953.31
612  1965-12-01  955.19
613  1966-01-01  985.93
614  1966-02-01  977.15
616  1966-04-01  943.70
646  1968-10-01  955.47
647  1968-11-01  964.12
648  1968-12-01  965.39
```

```
[110]: len(df1[df1['Price'] > upper_limit])
```

```
[110]: 9
```

```
[111]: df1.shape
```

```
[111]: (649, 2)
```

```
[112]: df2 =df1[(df1['Price'] > lower_limit) & (df1['Price'] < upper_limit)]
```

```
[113]: df2.shape
```

```
[113]: (640, 2)
```

```
[114]: df2.Price.value_counts()
```

```
[114]: Price
98.50      2
97.00      2
99.45      2
75.70      2
90.65      2
..
91.65      1
95.45      1
99.05      1
102.85     1
922.80     1
Name: count, Length: 628, dtype: int64
```

```
[115]: df2.columns
[115]: Index(['Date', 'Price'], dtype='object')

[116]: X =df2.drop(columns = ['Price'])

[117]: X.columns
[117]: Index(['Date'], dtype='object')

[118]: X =pd.get_dummies(X , drop_first = True).astype('int')

[119]: X.head()
[119]:
      Date
0 -17383680000000000000
1 -17356896000000000000
2 -17330112000000000000
3 -17305920000000000000
4 -17279136000000000000

[120]: X.shape
[120]: (640, 1)

[121]: y = df2[['Price']]
[122]: y.columns
[122]: Index(['Price'], dtype='object')

[123]: y.shape
[123]: (640, 1)

[124]: from sklearn.model_selection import train_test_split as split
[125]: X_train , X_test , y_train , y_test = split(X,y , train_size = 0.8 ,random_state = 42)
[126]: X_train.shape
[126]: (512, 1)

[127]: X_test.shape
[127]: (128, 1)

[128]: y_test.shape
```

```
[128]: (128, 1)
```

```
[129]: y_train.shape
```

```
[129]: (512, 1)
```

```
[130]: X.head()
```

```
[130]:          Date
0 -17383680000000000000
1 -17356896000000000000
2 -17330112000000000000
3 -17305920000000000000
4 -17279136000000000000
```

```
[131]: y.head()
```

```
[131]:      Price
0  55.00
1  56.55
2  56.00
3  58.30
4  66.45
```

```
[132]: X.columns
```

```
[132]: Index(['Date'], dtype='object')
```

```
[133]: y.columns
```

```
[133]: Index(['Price'], dtype='object')
```

```
[134]: X_test.head()
```

```
[134]:          Date
570 -23941440000000000000
265 -10413792000000000000
291 -97320960000000000000
597 -16830720000000000000
174 -12807936000000000000
```

```
[135]: y_test.head()
```

```
[135]:      Price
570  572.64
265  182.30
291  142.05
597  863.55
```

```
174 316.45
```

```
[136]: from sklearn.preprocessing import StandardScaler
[137]: sc = StandardScaler()
[138]: X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
[139]: X_train[:10]
[139]: array([[ 0.46501553],
[-0.40327054],
[ 0.57406695],
[ 0.59038879],
[ 0.79808871],
[-0.77974906],
[ 0.04997443],
[ 0.86857754],
[ 0.87951855],
[ 1.50225957]])
```

```
[140]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
[141]: df2 = df2.sort_values('Date').reset_index(drop=True)
[142]: df2['Pct_Change'] = df2['Price'].pct_change() * 100
[143]: df2['Next_Pct'] = df2['Pct_Change'].shift(-1)
df2['Target'] = (df2['Next_Pct'] > 0).astype(int)
[144]: df2 = df2.dropna()
[145]: X = df2[['Pct_Change']]
y = df2['Target']
[146]: X_train, X_test , y_train ,y_test = train_test_split(X , y , test_size = 0.2 ,random_state = 42)
[147]: from sklearn.linear_model import LogisticRegression
[148]: model = LogisticRegression()
[149]: model.fit(X_train , y_train)
[149]: LogisticRegression()
```

```
[150]: y_pred = model.predict(X_test)

[151]: df2 = df2.copy()

[152]: df2['MA_3'] = df2['Price'].rolling(3).mean()
df2['MA_7'] = df2['Price'].rolling(7).mean()
df2['Vol_5'] = df2['Pct_Change'].rolling(5).std()

df2 = df2.dropna()

X = df2[['Pct_Change', 'MA_3', 'MA_7', 'Vol_5']]
y = df2['Target']

[153]: df2['Return_1'] = df2['Price'].pct_change(1)
df2['Return_5'] = df2['Price'].pct_change(5)

# Trend indicators
df2['MA_10'] = df2['Price'].rolling(10).mean()
df2['MA_20'] = df2['Price'].rolling(20).mean()
df2['MA_Ratio'] = df2['MA_10'] / df2['MA_20']

df2['Vol_10'] = df2['Return_1'].rolling(10).std()

# Lag feature
df2['Lag_1'] = df2['Return_1'].shift(1)

[154]: features = ['Return_1', 'Return_5', 'MA_10', 'MA_20', 'MA_Ratio', 'Vol_10', ↴
    'Lag_1']
X = df2[features].dropna()
y = df2['Target'].loc[X.index]

[155]: accuracy = accuracy_score(y_test, y_pred)
print("Model Accuracy:", round(accuracy * 100, 2), "%")

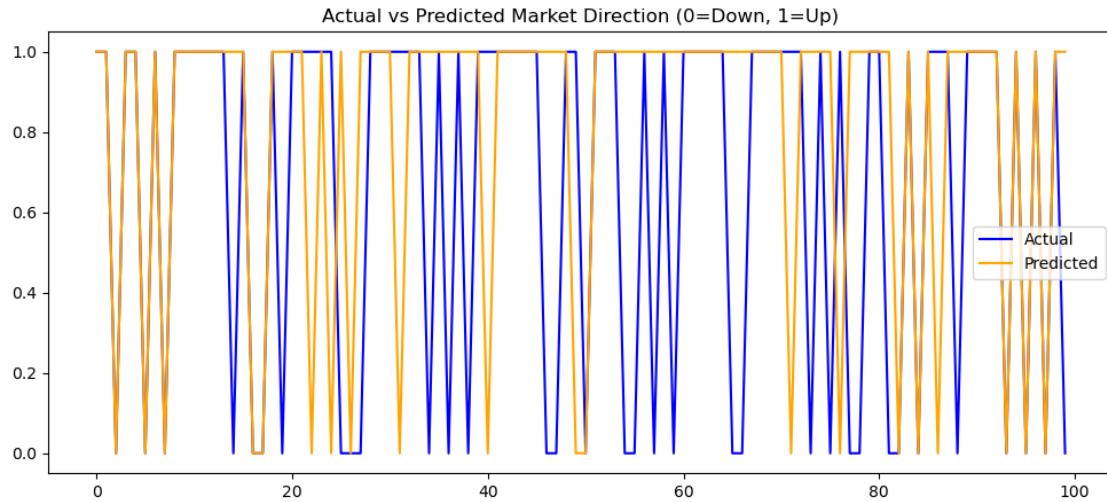
Model Accuracy: 68.75 %

[156]: df2['MA_3'] = df2['Price'].rolling(3).mean()
df2['MA_7'] = df2['Price'].rolling(7).mean()
df2['Vol_5'] = df2['Pct_Change'].rolling(5).std()

[157]: X = df2[['Pct_Change', 'MA_3', 'MA_7', 'Vol_5']].dropna()
y = df2['Target'].loc[X.index]

[158]: plt.figure(figsize=(12,5))
plt.plot(y_test.values[:100], label='Actual', color='blue')
plt.plot(y_pred[:100], label='Predicted', color='orange')
plt.title('Actual vs Predicted Market Direction (0=Down, 1=Up)')
```

```
plt.legend()  
plt.show()
```



```
[162]: from sklearn.model_selection import train_test_split  
from sklearn.preprocessing import StandardScaler  
from sklearn.pipeline import make_pipeline  
from sklearn.linear_model import Ridge  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
```

```
[163]: dfc = df2[['Date', 'Price']].copy()  
dfc['Date'] = pd.to_datetime(dfc['Date'], errors='coerce')  
dfc = dfc.dropna(subset=['Date', 'Price']).sort_values('Date').  
    ↪reset_index(drop=True)  
  
dfc['ret_1'] = dfc['Price'].pct_change(1)  
dfc['ma_5'] = dfc['Price'].rolling(5).mean()  
dfc['ma_10'] = dfc['Price'].rolling(10).mean()  
dfc['vol_5'] = dfc['ret_1'].rolling(5).std()  
dfc['lag_price_1'] = dfc['Price'].shift(1)  
dfc['target_next_price'] = dfc['Price'].shift(-1)  
  
dfc = dfc.  
    ↪dropna(subset=['ret_1', 'ma_5', 'ma_10', 'vol_5', 'lag_price_1', 'target_next_price']).  
    ↪reset_index(drop=True)
```

```
[164]: X = dfc[['lag_price_1', 'ret_1', 'ma_5', 'ma_10', 'vol_5']]
y = dfc['target_next_price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, ↳
shuffle=False)

linreg = make_pipeline(StandardScaler(), Ridge(alpha=1.0))
linreg.fit(X_train, y_train)
y_pred_lr = linreg.predict(X_test)

rf = RandomForestRegressor(n_estimators=300, max_depth=8, random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
```

```
[165]: def report(name, y_true, y_pred):
    r2 = r2_score(y_true, y_pred)
    mae = mean_absolute_error(y_true, y_pred)
    print(f"{name:20s} | R²: {r2:.4f} | MAE: {mae:.4f}")
    return r2, mae

r2_lr, mae_lr = report("Linear Regression", y_test, y_pred_lr)
r2_rf, mae_rf = report("Random Forest", y_test, y_pred_rf)
```

Linear Regression | R²: 0.9590 | MAE: 21.7342
 Random Forest | R²: -2.8758 | MAE: 239.6037

```
[166]: plt.figure(figsize=(12,5))
plt.plot(y_test.values, label="Actual Price", color='black', linewidth=2)
plt.plot(y_pred_lr, label="Linear Regression", alpha=0.7)
plt.plot(y_pred_rf, label="Random Forest", alpha=0.7)
plt.title("Actual vs Predicted Prices (Comparison)")
plt.xlabel("Time")
plt.ylabel("Price")
plt.legend()
plt.tight_layout()
plt.show()

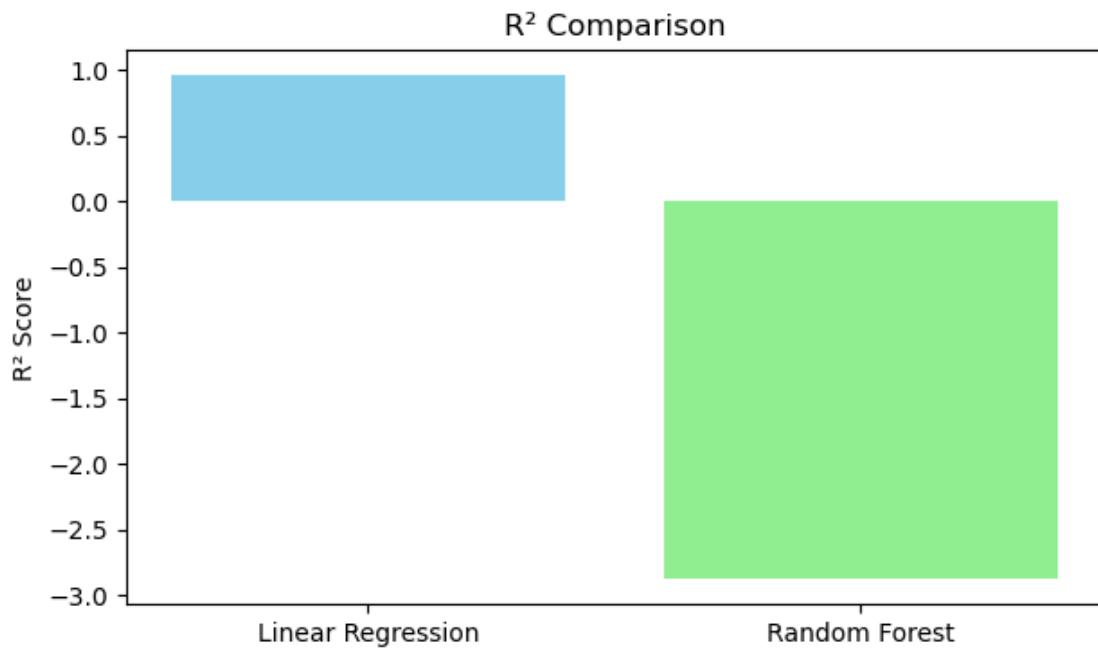
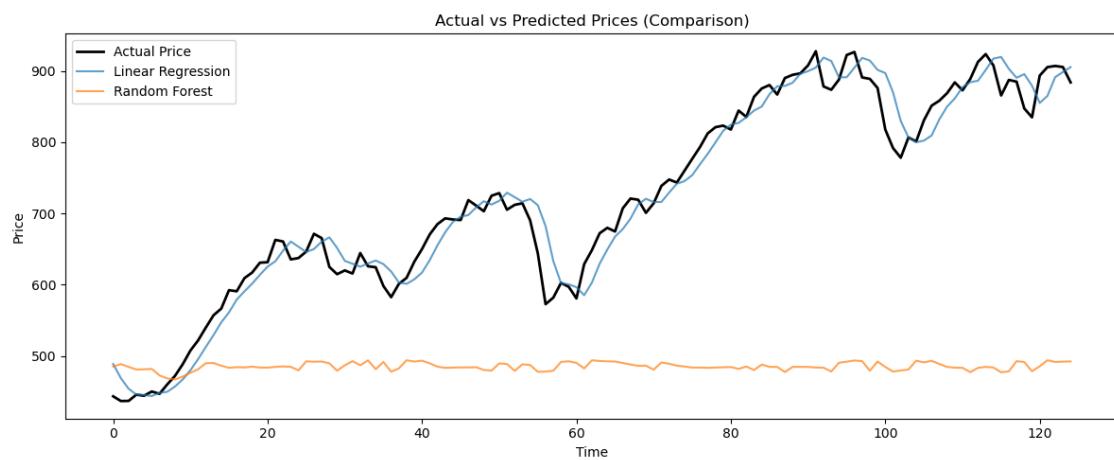
models = ["Linear Regression", "Random Forest"]
r2_values = [r2_lr, r2_rf]
mae_values = [mae_lr, mae_rf]
```

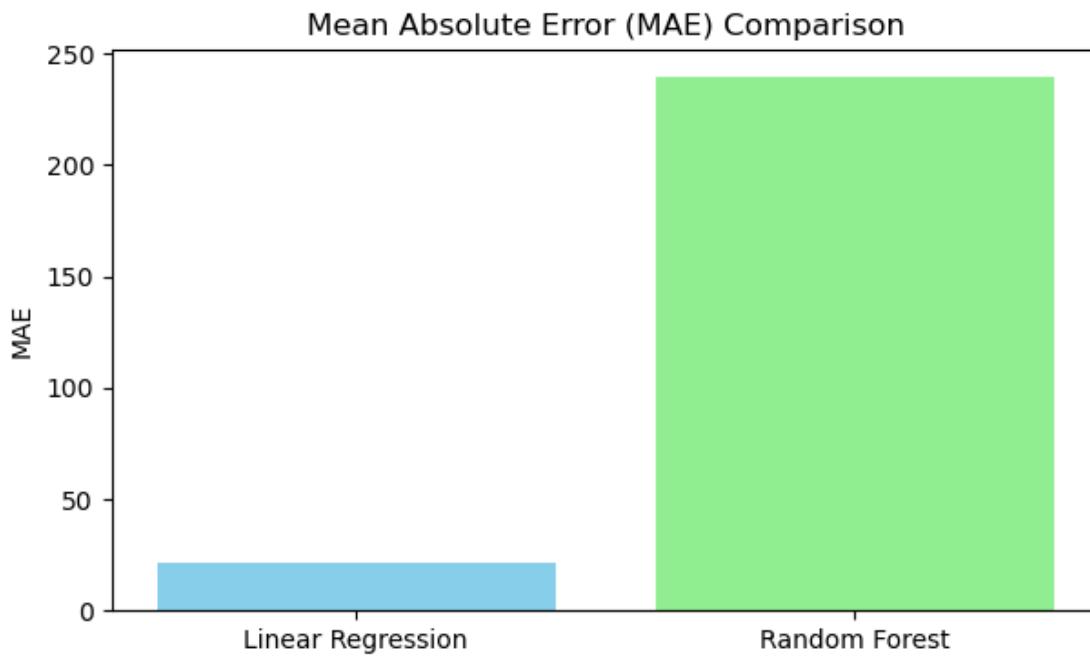
```

fig, ax1 = plt.subplots(1, 1, figsize=(7,4))
ax1.bar(models, r2_values, color=['skyblue', 'lightgreen'])
ax1.set_title("R2 Comparison")
ax1.set_ylabel("R2 Score")
plt.show()

fig, ax2 = plt.subplots(1, 1, figsize=(7,4))
ax2.bar(models, mae_values, color=['skyblue', 'lightgreen'])
ax2.set_title("Mean Absolute Error (MAE) Comparison")
ax2.set_ylabel("MAE")
plt.show()

```





```
[167]: from sklearn import datasets
```

```
[168]: dfdata = sns.load_dataset("dowjones")
dfdata.head()
```

```
[168]:      Date  Price
0 1914-12-01  55.00
1 1915-01-01  56.55
2 1915-02-01  56.00
3 1915-03-01  58.30
4 1915-04-01  66.45
```

```
[169]: df = pd.DataFrame(dfdata.Date , columns = dfdata.Price)
```

```
[170]: df.head()
```

```
[170]: Empty DataFrame
Columns: [55.0, 56.55, 56.0, 58.3, 66.45, 65.95, 68.4, 71.85, 79.25, 85.5,
92.35, 94.35, 97.0, 94.7, 93.55, 93.3, 89.75, 90.15, 90.65, 88.45, 91.0, 97.45,
102.1, 107.9, 98.5, 97.15, 90.95, 94.65, 93.9, 93.35, 96.95, 92.9, 88.65, 83.6,
79.05, 71.4, 70.2, 76.6, 79.95, 78.05, 77.65, 81.05, 80.45, 81.85, 81.95, 82.5,
86.25, 84.0, 82.5, 81.65, 82.45, 86.55, 91.15, 99.4, 103.6, 109.7, 103.25,
108.2, 113.9, 111.6, 105.8, 105.9, 96.5, 97.95, 99.45, 91.1, 91.7, 90.7, 85.25,
```

```
86.5, 84.85, 79.3, 72.2, 74.75, 75.7, 75.05, 77.0, 76.7, 69.2, 68.6, 66.95,  
69.35, 71.7, 75.7, 79.8, 80.6, 83.75, 87.2, 91.3, 93.95, 93.55, 95.0, 98.5,  
99.2, 99.75, 95.75, 97.0, 98.2, 100.8, 103.9, ...]  
Index: []
```

```
[0 rows x 649 columns]
```

```
[171]: df['Price'] = dfdata['Price']
```

```
[172]: from sklearn.model_selection import train_test_split as split
```

```
[173]: X_train , X_test , y_train , y_test = split(dfdata.Price , dfdata.Price ,  
    ↪test_size = 0.2)
```

```
[174]: from sklearn import svm
```

```
[179]: import numpy as np
```

```
if not isinstance(X_train, np.ndarray):  
    X_train = X_train.to_numpy()  
if not isinstance(X_test, np.ndarray):  
    X_test = X_test.to_numpy()  
if not isinstance(y_train, np.ndarray):  
    y_train = np.ravel(y_train)  
if not isinstance(y_test, np.ndarray):  
    y_test = np.ravel(y_test)  
  
if X_train.ndim == 1:  
    X_train = X_train.reshape(-1, 1)  
if X_test.ndim == 1:  
    X_test = X_test.reshape(-1, 1)
```

```
[180]: X_train = np.asarray(X_train)  
X_test = np.asarray(X_test)  
y_train = np.asarray(y_train)  
y_test = np.asarray(y_test)
```

```
if X_train.ndim == 1:  
    X_train = X_train.reshape(-1, 1)  
if X_test.ndim == 1:  
    X_test = X_test.reshape(-1, 1)  
y_train = y_train.flatten()  
y_test = y_test.flatten()
```

```
[181]: from sklearn.model_selection import GridSearchCV
```

```
[182]: srv = GridSearchCV(svm.SVC(gamma = 'auto'),
                         {'C' :[1,10,20],
                          'kernel':['rbf' , 'linear']},
                         cv = 5,return_train_score = False)

[183]: X = dfc[['ret_1','ma_5','vol_5']]
y = dfc['Price']

[184]: from sklearn.svm import SVR
pipe = make_pipeline(StandardScaler(), SVR())

param_grid = {'svr__kernel': ['rbf', 'linear'],
              'svr__C': [0.1, 1, 10, 100],
              'svr__gamma': ['scale', 'auto']}

grid = GridSearchCV(pipe, param_grid, scoring='neg_mean_squared_error', cv=3,
                     n_jobs=-1, verbose=0)
grid.fit(X_train, y_train)

[184]: GridSearchCV(cv=3,
                    estimator=Pipeline(steps=[('standardscaler', StandardScaler()),
                                              ('svr', SVR())]),
                    n_jobs=-1,
                    param_grid={'svr__C': [0.1, 1, 10, 100],
                                'svr__gamma': ['scale', 'auto'],
                                'svr__kernel': ['rbf', 'linear']},
                    scoring='neg_mean_squared_error')

[185]: param_grid = {
    'kernel': ['linear', 'rbf'],
    'C': [0.1, 1, 10],
    'gamma': ['scale', 'auto']
}

grid_search = GridSearchCV(SVR(), param_grid, cv=5)
grid_search.fit(X_train, y_train)

[185]: GridSearchCV(cv=5, estimator=SVR(),
                    param_grid={'C': [0.1, 1, 10], 'gamma': ['scale', 'auto'],
                                'kernel': ['linear', 'rbf']}))

[186]: srv = GridSearchCV(model, param_grid, cv=3, scoring='r2', n_jobs=-1)

[189]: from sklearn.svm import SVR

svr = SVR()
svr.fit(X_train, y_train)
```

```

y_pred_svr = svr.predict(X_test)

[190]: param_grid = {
    'kernel': ['linear', 'rbf'],
    'C': [0.1, 1, 10, 100],
    'gamma': ['scale', 'auto']
}

[191]: grid_search.cv_results_

[191]: {'mean_fit_time': array([0.00108385, 0.00828385, 0.0006237 , 0.00573945,
 0.0005024 ,
 0.00542922, 0.00045662, 0.00534458, 0.00071254, 0.00850444,
 0.00064778, 0.0095192 ]),
 'std_fit_time': array([3.41996472e-04, 1.75079654e-03, 1.02461754e-04,
 4.07793719e-04,
 2.00188710e-05, 1.34500320e-04, 2.32989486e-05, 3.63514655e-04,
 1.90951856e-04, 9.57246140e-04, 1.38559230e-04, 1.50043790e-03]),
 'mean_score_time': array([0.00085921, 0.00567489, 0.00061202, 0.0043458 ,
 0.00049405,
 0.0040956 , 0.00051703, 0.00411634, 0.00066457, 0.00485358,
 0.00068817, 0.00463767]),
 'std_score_time': array([1.34080080e-04, 1.00551961e-03, 6.66561383e-05,
 2.93400643e-04,
 5.62465761e-06, 1.43861652e-04, 7.74348310e-05, 3.37990150e-04,
 1.31205156e-04, 4.34260409e-04, 9.34912912e-05, 1.88398319e-04]),
 'param_C': masked_array(data=[0.1, 0.1, 0.1, 0.1, 1.0, 1.0, 1.0, 1.0, 10.0,
 10.0,
 10.0, 10.0],
 mask=[False, False, False, False, False, False, False, False,
 False, False, False],
 fill_value=1e+20),
 'param_gamma': masked_array(data=['scale', 'scale', 'auto', 'auto', 'scale',
 'scale',
 'auto', 'auto', 'scale', 'scale', 'auto', 'auto'],
 mask=[False, False, False, False, False, False, False, False,
 False, False, False],
 fill_value='?',
 dtype=object),
 'param_kernel': masked_array(data=['linear', 'rbf', 'linear', 'rbf', 'linear',
 'rbf',
 'linear', 'rbf', 'linear', 'rbf', 'linear', 'rbf'],
 mask=[False, False, False, False, False, False, False, False,
 False, False, False],
 fill_value='?',
 dtype=object),
 'params': [{'C': 0.1, 'gamma': 'scale', 'kernel': 'linear'}],
}
```

```

{'C': 0.1, 'gamma': 'scale', 'kernel': 'rbf'},
{'C': 0.1, 'gamma': 'auto', 'kernel': 'linear'},
{'C': 0.1, 'gamma': 'auto', 'kernel': 'rbf'},
{'C': 1, 'gamma': 'scale', 'kernel': 'linear'},
{'C': 1, 'gamma': 'scale', 'kernel': 'rbf'},
{'C': 1, 'gamma': 'auto', 'kernel': 'linear'},
{'C': 1, 'gamma': 'auto', 'kernel': 'rbf'},
{'C': 10, 'gamma': 'scale', 'kernel': 'linear'},
{'C': 10, 'gamma': 'scale', 'kernel': 'rbf'},
{'C': 10, 'gamma': 'auto', 'kernel': 'linear'},
{'C': 10, 'gamma': 'auto', 'kernel': 'rbf']},
'split0_test_score': array([ 0.99999993, -0.27503558,  0.99999993, -0.34674163,
0.99999993,
     0.27957281,  0.99999993, -0.33860568,  0.99999993,  0.98978917,
0.99999993, -0.26658791]),
'split1_test_score': array([ 0.99999993, -0.22531024,  0.99999993, -0.30073595,
0.99999993,
     0.33822966,  0.99999993, -0.30149746,  0.99999993,  0.99187353,
0.99999993, -0.252192 ]),
'split2_test_score': array([ 0.99999992, -0.09035798,  0.99999992, -0.15250005,
0.99999992,
     0.37965392,  0.99999992, -0.14892305,  0.99999992,  0.98028337,
0.99999992, -0.09434162]),
'split3_test_score': array([ 0.99999987,  0.0057193 ,  0.99999987, -0.06809463,
0.99999987,
     0.54136339,  0.99999987, -0.06465951,  0.99999987,  0.99692366,
0.99999987, -0.01147313]),
'split4_test_score': array([ 0.99999993, -0.20533687,  0.99999993, -0.27301928,
0.99999993,
     0.30530087,  0.99999993, -0.27083871,  0.99999993,  0.98153106,
0.99999993, -0.22806874]),
'mean_test_score': array([ 0.99999991, -0.15806427,  0.99999991, -0.22821831,
0.99999991,
     0.36882413,  0.99999991, -0.22490488,  0.99999991,  0.98808016,
0.99999991, -0.17053268]),
'std_test_score': array([2.44190805e-08, 1.01831436e-01, 2.44190805e-08,
1.02651409e-01,
     2.44190805e-08, 9.25463174e-02, 2.44190805e-08, 1.02328910e-01,
2.44190805e-08, 6.31186984e-03, 2.44190805e-08, 1.00309786e-01]),
'rank_test_score': array([ 1,  9,  1, 12,  1,  8,  1, 11,  1,  7,  1, 10],
dtype=int32)}

```

[192]: `srv = GridSearchCV(model, param_grid, cv=3, scoring='r2', n_jobs=-1)`

```
srv.fit(X_train, y_train)
```

```
best = pd.DataFrame(srv.cv_results_)
print(srv.best_params_)

{'C': 0.1, 'gamma': 'scale', 'kernel': 'linear'}
```

```
[193]: srv.best_score_
```

```
[193]: 0.9999999194621404
```