

S&P_500_Dataset

By Hatem Elgenedy

November 7, 2025

```
[141]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
from datetime import datetime
from dateutil.relativedelta import relativedelta
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
```

```
[142]: import os

file_path = "/Users/hatemelgenedy/Desktop/AI and Data Science Microsoft course/
↳Projects/S& P 500 dataset/sp500_index.csv"

if not os.path.exists(file_path):
    raise FileNotFoundError(f"CSV file not found at {file_path}")

# read the CSV; skip malformed lines if any
df = pd.read_csv(file_path, on_bad_lines='skip')
df.head(6)
```

```
[142]:      Date   S&P500
0  2014-12-22  2078.54
1  2014-12-23  2082.17
2  2014-12-24  2081.88
3  2014-12-26  2088.77
4  2014-12-29  2090.57
5  2014-12-30  2080.35
```

```
[143]: !pip install ydata-profiling
from ydata_profiling import ProfileReport
profile = ProfileReport(df, title="S&P 500 Dataset Profiling Report",
↳explorative=True)
profile.to_notebook_iframe()
```

```
Requirement already satisfied: ydata-profiling in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (4.17.0)
Requirement already satisfied: scipy<1.16,>=1.4.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (1.13.1)
Requirement already satisfied: pandas!=1.4.0,<3.0,>1.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (2.3.1)
Requirement already satisfied: matplotlib<=3.10,>=3.5 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (3.10.0)
Requirement already satisfied: pydantic>=2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (2.11.7)
Requirement already satisfied: PyYAML<6.1,>=5.0.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (6.0.2)
Requirement already satisfied: jinja2<3.2,>=2.11.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (3.1.6)
Requirement already satisfied: visions<0.8.2,>=0.7.5 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling) (0.8.1)
Requirement already satisfied: numpy<2.2,>=1.16.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (1.26.4)
Requirement already satisfied: minify-html>=0.15.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (0.18.1)
Requirement already satisfied: filetype>=1.0.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (1.2.0)
Requirement already satisfied: phik<0.13,>=0.11.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (0.12.5)
Requirement already satisfied: requests<3,>=2.24.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (2.32.4)
Requirement already satisfied: tqdm<5,>=4.48.2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (4.67.1)
Requirement already satisfied: seaborn<0.14,>=0.10.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (0.13.2)
Requirement already satisfied: multimethod<2,>=1.4 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (1.12)
Requirement already satisfied: statsmodels<1,>=0.13.2 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-profiling) (0.14.5)
Requirement already satisfied: typeguard<5,>=3 in /opt/anaconda3/envs/anaconda-
```

```
nlp/lib/python3.11/site-packages (from ydata-profiling) (4.4.4)
Requirement already satisfied: imagehash==4.3.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (4.3.1)
Requirement already satisfied: wordcloud>=1.9.3 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (1.9.4)
Requirement already satisfied: dacite>=1.8 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from ydata-profiling) (1.9.2)
Requirement already satisfied: numba<=0.61,>=0.56.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from ydata-
profiling) (0.61.0)
Requirement already satisfied: PyWavelets in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from imagehash==4.3.1->ydata-profiling)
(1.9.0)
Requirement already satisfied: pillow in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from imagehash==4.3.1->ydata-profiling)
(11.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from jinja2<3.2,>=2.11.1->ydata-profiling)
(3.0.2)
Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
(1.3.1)
Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
(0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
matplotlib<=3.10,>=3.5->ydata-profiling) (4.55.3)
Requirement already satisfied: kiwisolver>=1.3.1 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
matplotlib<=3.10,>=3.5->ydata-profiling) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
(24.2)
Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from matplotlib<=3.10,>=3.5->ydata-profiling)
(3.2.0)
Requirement already satisfied: python-dateutil>=2.7 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
matplotlib<=3.10,>=3.5->ydata-profiling) (2.9.0.post0)
Requirement already satisfied: llvmlite<0.45,>=0.44.0dev0 in
/opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from
numba<=0.61,>=0.56.0->ydata-profiling) (0.44.0)
Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling)
(2025.2)
Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/envs/anaconda-
nlp/lib/python3.11/site-packages (from pandas!=1.4.0,<3.0,>1.1->ydata-profiling)
```

(2025.2)

Requirement already satisfied: joblib>=0.14.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from phik<0.13,>=0.11.1->ydata-profiling)
(1.5.1)

Requirement already satisfied: charset_normalizer<4,>=2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (3.3.2)

Requirement already satisfied: idna<4,>=2.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling)
(3.7)

Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (2.5.0)

Requirement already satisfied: certifi>=2017.4.17 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from requests<3,>=2.24.0->ydata-profiling) (2025.8.3)

Requirement already satisfied: patsy>=0.5.6 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from statsmodels<1,>=0.13.2->ydata-profiling)
(1.0.2)

Requirement already satisfied: typing_extensions>=4.14.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from typeguard<5,>=3->ydata-profiling) (4.15.0)

Requirement already satisfied: attrs>=19.3.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
(24.3.0)

Requirement already satisfied: networkx>=2.4 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
(3.4.2)

Requirement already satisfied: puremagic in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from visions<0.8.2,>=0.7.5->visions[type_image_path]<0.8.2,>=0.7.5->ydata-profiling)
(1.30)

Requirement already satisfied: annotated-types>=0.6.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pydantic>=2->ydata-profiling) (0.6.0)

Requirement already satisfied: pydantic-core==2.33.2 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pydantic>=2->ydata-profiling) (2.33.2)

Requirement already satisfied: typing-inspection>=0.4.0 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from pydantic>=2->ydata-profiling) (0.4.0)

Requirement already satisfied: six>=1.5 in /opt/anaconda3/envs/anaconda-nlp/lib/python3.11/site-packages (from python-dateutil>=2.7->matplotlib<=3.10,>=3.5->ydata-profiling) (1.17.0)

Summarize dataset: 0% | 0/5 [00:00<?, ?it/s]

```

100%|      | 2/2 [00:00<00:00, 12.46it/s]
Generate report structure:  0%|          | 0/1 [00:00<?, ?it/s]
Render HTML:   0%|          | 0/1 [00:00<?, ?it/s]
<IPython.core.display.HTML object>

[144]: profile.to_file("S&P_500_report.html")

Export report to file:  0%|          | 0/1 [00:00<?, ?it/s]

[145]: df.isnull().sum()

[145]: Date      0
       S&P500    0
       dtype: int64

[146]: duplicates = df . duplicated().sum()
       print(duplicates)

0

[147]: df.shape

[147]: (2517, 2)

[148]: print(df.columns)

Index(['Date', 'S&P500'], dtype='object')

[149]: companies_file_path = "/Users/hatemelgenedy/Desktop/AI and Data Science\u2192Microsoft course/Projects/S& P 500 dataset/sp500_companies.csv"


df = pd.read_csv(companies_file_path, on_bad_lines='skip')

df[["Weight"]] = pd.to_numeric(df[["Weight"]], errors="coerce")
df[["Marketcap"]] = pd.to_numeric(df[["Marketcap"]], errors="coerce")
df = df.dropna(subset=["Weight", "Marketcap"]).reset_index(drop=True)

[150]: plt.figure(figsize = (4,2))
       sns.heatmap(df.isnull())

[150]: <Axes: >

```



```
[151]: null_percentage = df.isnull()/df.shape[0]*100
null_percentage
```

```
[151]:      Exchange   Symbol   Shortname   Longname   Sector   Industry   Currentprice \
0          0.0       0.0       0.0       0.0       0.0       0.0       0.0
1          0.0       0.0       0.0       0.0       0.0       0.0       0.0
2          0.0       0.0       0.0       0.0       0.0       0.0       0.0
3          0.0       0.0       0.0       0.0       0.0       0.0       0.0
4          0.0       0.0       0.0       0.0       0.0       0.0       0.0
..
497        0.0       0.0       0.0       0.0       0.0       0.0       0.0
498        0.0       0.0       0.0       0.0       0.0       0.0       0.0
499        0.0       0.0       0.0       0.0       0.0       0.0       0.0
500        0.0       0.0       0.0       0.0       0.0       0.0       0.0
501        0.0       0.0       0.0       0.0       0.0       0.0       0.0

      Marketcap   Ebitda   Revenuegrowth   City   State   Country \
0          0.0       0.0       0.0       0.0       0.0       0.0
1          0.0       0.0       0.0       0.0       0.0       0.0
2          0.0       0.0       0.0       0.0       0.0       0.0
3          0.0       0.0       0.0       0.0       0.0       0.0
4          0.0       0.0       0.0       0.0       0.0       0.0
..
497        0.0       0.0       0.0       0.0       0.0       0.0
498        0.0       0.0       0.0       0.0       0.0       0.0
```

```

499      0.0    0.0      0.0    0.0    0.0    0.0
500      0.0    0.0      0.0    0.0    0.0    0.0
501      0.0    0.0      0.0    0.0    0.0    0.0

```

	Fulltimeemployees	Longbusinesssummary	Weight
0	0.000000	0.0	0.0
1	0.000000	0.0	0.0
2	0.000000	0.0	0.0
3	0.000000	0.0	0.0
4	0.000000	0.0	0.0
..
497	0.000000	0.0	0.0
498	0.000000	0.0	0.0
499	0.000000	0.0	0.0
500	0.000000	0.0	0.0
501	0.199203	0.0	0.0

[502 rows x 16 columns]

[152]: df.shape

[152]: (502, 16)

```

[153]: def iqr_bounds(s, k=1.5):
    q1,q3 = s.quantile(0.25) , s.quantile(0.75)
    iqr = q3 - q1
    return q1 - k*iqr , q3 +k*iqr

df["log10_mcap"] = np.log10(df["Marketcap"].astype(float))

w_low, w_high = iqr_bounds(df["Weight"])
m_low , m_high = iqr_bounds(df["log10_mcap"])

mask_w = (df["Weight"] <w_low) | (df ["Weight"] > w_high)
mask_m = (df["log10_mcap"] < m_low) | (df["log10_mcap"] > m_high)
mask_any = mask_w | mask_m

df["Outlier_Weight"] = mask_w
df["Outlier_Mcaplog"] = mask_m
df["Outlier_Any"] = mask_any

label_col = "Shortname" if "Shortname" in df.columns else ("Longname" if
    ↵"Longname" in df.columns else "Symbol")
outliers = df.loc[mask_any , ["Symbol" , label_col , "Weight" , "Marketcap"]]

```

```
print("IQR thresholds :")
print(f" Weight : low = {w_low : .3f} , high = {w_high : .3f}")
print("\nOutier countd : \n" , df[["Outlier_Weight" , "Outlier_Mcaplog" , ↴"Outlier_Any"]])
```

IQR thresholds :
Weight : low = -0.001 , high = 0.003

Outier countd :


	Outlier_Weight	Outlier_Mcaplog	Outlier_Any
0	True	True	True
1	True	True	True
2	True	True	True
3	True	True	True
4	True	True	True
..
497	False	False	False
498	False	False	False
499	False	False	False
500	False	False	False
501	False	False	False

[502 rows x 3 columns]

```
[154]: import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(8,5))
sns.boxplot(x=df["Weight"])
plt.title("S&P 500 Company Weights (with Outliers)")
plt.xlabel("Weight")
plt.show()
```

S&P 500 Company Weights (with Outliers)



```
[185]: plt.figure(figsize = (10,7))
sns.stripplot(data = df , x = 'Sector' , y = 'Weight' , color = 'r' , jitter = True , alpha = 0.5)

plt.xticks(rotation = 45 , ha = 'right')
plt.tight_layout()
```



```
[156]: model = LinearRegression()

[157]: x = df[['Sector']]
y = df[['Weight']]

[158]: X = df.drop(columns = ['Sector'])
y = df[['Weight']]

[159]: df_clean = df.dropna(subset = ["Weight", "Sector"]).reset_index(drop = True)
X = df_clean[["Sector"]].to_numpy()
y = df_clean[["Weight"]].to_numpy()

[160]: X = pd.get_dummies(df[['Sector']], drop_first=True)
y = df[['Weight']].apply(pd.to_numeric, errors='coerce')
mask = y['Weight'].notna()
X, y = X.loc[mask], y.loc[mask]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

model = LinearRegression().fit(X_train, y_train)
```

```
[161]: X = pd.get_dummies(df[['Sector']], drop_first = True)
y = pd.to_numeric(df['Weight'], errors = 'coerce')

[162]: X_train , X_test , y_train , y_test = train_test_split(X, y , test_size = 0.2 ,random_state = 42)

[163]: print('80% input traning data :' , X_train.shape)
print('20% input testing data :' , X_test.shape)

80% input traning data : (401, 10)
20% input testing data : (101, 10)

[164]: print("X_train shape :" , X_train.shape)
print("X_test shape :" , X_test.shape)
print("y_train shape :" , y_train.shape)
print("y_test shape" , y_test.shape)

X_train shape : (401, 10)
X_test shape : (101, 10)
y_train shape : (401,)
y_test shape (101,)

[165]: from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import OneHotEncoder ,StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error

[166]: model = LinearRegression()

[167]: model.fit(X_train , y_train)

[167]: LinearRegression()

[168]: y_pred = model.predict(X_test)

[186]: features = ['Marketcap', 'RevenueGrowth', 'Employees', 'PEratio', 'DividendYield']
target = 'Weight'

[187]: y_pred[:5]

[187]: array([0.0023986 , 0.0017719 , 0.00237563, 0.00067109, 0.00237563])

[170]: y_test[:5]

[170]: 268      0.000577
73       0.002450
289      0.000532
```

```
155    0.001159  
104    0.001605  
Name: Weight, dtype: float64
```

```
[189]: y_pred_1d = np.asarray(y_pred).ravel()  
y_test_1d = np.asarray(y_test).ravel()
```

```
[190]: compare = pd.DataFrame({  
    "Predicted" : np.ravel(y_pred),  
    "Actual": np.ravel(y_test) if not hasattr(y_test, "to_numpy") else  
    ↪y_test.to_numpy().ravel()})  
print(compare.head())
```

	Predicted	Actual
0	0.002399	0.000577
1	0.001772	0.002450
2	0.002376	0.000532
3	0.000671	0.001159
4	0.002376	0.001605

```
[195]: from sklearn.metrics import mean_squared_error
```

```
[196]: r2 = r2_score(y_test, y_pred)  
mae = mean_absolute_error(y_test, y_pred)  
rmse = np.sqrt(mean_squared_error(y_test, y_pred))  
  
print("R²:", r2)  
print("MAE:", mae)  
print("RMSE:", rmse)
```

R²: -0.016411245199545066
MAE: 0.0027731236401882017
RMSE: 0.009123135476645521

```
[191]: compare = pd.DataFrame({"Predicted" :np.ravel(y_pred),  
                            "Actual": np.ravel(y_test)}, index = X_test.index)
```

```
[192]: from sklearn.metrics import r2_score
```

```
[193]: print(r2_score(y_pred,y_test))
```

-36.41707159049202

```
[176]: from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error  
import numpy as np  
  
y_true = np.ravel(y_test)  
y_hat = np.ravel(y_pred)
```

```

r2 = r2_score(y_true, y_hat)
mae = mean_absolute_error(y_true, y_hat)
rmse = np.sqrt(mean_squared_error(y_true, y_hat))
print(f"R²: {r2:.4f}")
print(f"MAE: {mae:.4f}")
print(f"RMSE: {rmse:.4f}")

```

R²: -0.0164
MAE: 0.0028
RMSE: 0.0091

```

[177]: y_true = np.ravel(y_test)
y_hat = np.ravel(y_pred)

# Baseline: always predict mean(y_train)
y_base = np.full_like(y_true, fill_value=np.mean(y_train))

print("Model vs Baseline")
print("R2 (model):", r2_score(y_true, y_hat))
print("R2 (base) : ", r2_score(y_true, y_base))           # should be ~0 by ↴definition
print("MAE (model):", mean_absolute_error(y_true, y_hat))
print("MAE (base) : ", mean_absolute_error(y_true, y_base))
print("RMSE(model):", np.sqrt(mean_squared_error(y_true, y_hat)))
print("RMSE(base) : ", np.sqrt(mean_squared_error(y_true, y_base)))

```

Model vs Baseline
R2 (model): -0.016411245199545066
R2 (base) : -0.010383930477205539
MAE (model): 0.0027731236401882017
MAE (base) : 0.0027448834996816134
RMSE(model): 0.009123135476645521
RMSE(base) : 0.009096045176677364

```

[178]: y_range = y_true.max() - y_true.min()
y_std = y_true.std()

nrmse_range = np.sqrt(mean_squared_error(y_true, y_hat)) / (y_range if y_range!=0 else 1)
nrmse_std = np.sqrt(mean_squared_error(y_true, y_hat)) / (y_std if y_std!=0 else 1)
print("NRMSE (fraction of range):", nrmse_range)
print("NRMSE (fraction of std) : ", nrmse_std)

```

NRMSE (fraction of range): 0.13204967848105265
NRMSE (fraction of std) : 1.0081722299287683

```
[198]: from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
rf = RandomForestRegressor(n_estimators=300, max_depth=10, random_state=42)
rf.fit(X_scaled, y)
y_pred = rf.predict(X_scaled)
```

```
[199]: r2 = r2_score(y, y_pred)
mae = mean_absolute_error(y, y_pred)
rmse = np.sqrt(mean_squared_error(y, y_pred))

print(f"R²: {r2:.3f}")
print(f"MAE: {mae:.4f}")
print(f"RMSE: {rmse:.4f}")
```

R²: 0.049
MAE: 0.0022
RMSE: 0.0060

```
[202]: print(df.columns)
```

Index(['Exchange', 'Symbol', 'Shortname', 'Longname', 'Sector', 'Industry',
'Currentprice', 'Marketcap', 'Ebitda', 'Revenuegrowth', 'City', 'State',
'Country', 'Fulltimeemployees', 'Longbusinesssummary', 'Weight',
'log10_mcap', 'Outlier_Weight', 'Outlier_Mcaplog', 'Outlier_Any'],
dtype='object')

```
[204]: features = ['Marketcap', 'Currentprice', 'Ebitda', 'Revenuegrowth',  

    ↪'Fulltimeemployees', 'log10_mcap']
target = 'Weight'

# Drop missing values
df_clean = df[features + [target]].dropna()

X = df_clean[features]
y = df_clean[target]

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  

    ↪random_state=42)

# Scale features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Gradient Boosting Model
```

```

gb = GradientBoostingRegressor(n_estimators=500, learning_rate=0.05,
    max_depth=5, random_state=42)
gb.fit(X_train_scaled, y_train)
y_pred = gb.predict(X_test_scaled)

# Evaluate performance
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))

print(f"R²: {r2:.3f}")
print(f"MAE: {mae:.6f}")
print(f"RMSE: {rmse:.6f}")


```

R²: 0.936
MAE: 0.000256
RMSE: 0.001877


```
[205]: from sklearn.model_selection import cross_val_score
cv_scores = cross_val_score(gb, X_train_scaled, y_train, cv=5, scoring='r2')
print("Cross-validated R²:", np.mean(cv_scores))
```

Cross-validated R²: 0.956206476977691

```
[206]: plt.figure(figsize=(8,6))
sns.scatterplot(x=y_test, y=y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.title("Actual vs Predicted S&P 500 Weights")
plt.xlabel("Actual Weight")
plt.ylabel("Predicted Weight")
plt.grid(True, linestyle='--', alpha=0.5)
plt.show()
```



```
[217]: importance = pd.Series(gb.feature_importances_, index=features).
    sort_values(ascending=True)
importance.plot(kind='barh', figsize=(8,4), title='Feature Importance (Gradient
    Boosting)')
plt.show()
```




```
[209]: features = [
    'Marketcap',
    'Currentprice',
    'Ebitda',
    'Revenuegrowth',
    'Fulltimeemployees',
    'log10_mcap'      # this is already a great feature
]
target = 'Weight'
```

```
[210]: from sklearn.ensemble import GradientBoostingRegressor

gb = GradientBoostingRegressor(n_estimators=500, learning_rate=0.05,
    max_depth=5, random_state=42)
gb.fit(X_train, y_train)
y_pred_gb = gb.predict(X_test)
```

```
[211]: plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred, alpha=0.7, color='blue', label='Predicted vs Actual')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--',
    label='Perfect Fit')

plt.title("Model Predictions vs Actual Weights")
plt.xlabel("Actual Weight")
plt.ylabel("Predicted Weight")
plt.legend()
plt.show()
```



```
[212]: df.head()
```

	Exchange	Symbol	Shortname	Longname \
0	NMS	AAPL	Apple Inc.	Apple Inc.
1	NMS	NVDA	NVIDIA Corporation	NVIDIA Corporation
2	NMS	MSFT	Microsoft Corporation	Microsoft Corporation
3	NMS	AMZN	Amazon.com, Inc.	Amazon.com, Inc.
4	NMS	GOOGL	Alphabet Inc.	Alphabet Inc.

	Sector	Industry	Currentprice \
0	Technology	Consumer Electronics	254.49
1	Technology	Semiconductors	134.70
2	Technology	Software - Infrastructure	436.60
3	Consumer Cyclical	Internet Retail	224.92
4	Communication Services	Internet Content & Information	191.41

```

      Marketcap      Ebitda Revenuegrowth          City State \
0  3846819807232  1.346610e+11       0.061    Cupertino   CA
1  3298803056640  6.118400e+10       1.224    Santa Clara  CA
2  3246068596736  1.365520e+11       0.160    Redmond    WA
3  2365033807872  1.115830e+11       0.110    Seattle    WA
4  2351625142272  1.234700e+11       0.151  Mountain View  CA

      Country Fulltimeemployees \
0  United States        164000.0
1  United States        29600.0
2  United States        228000.0
3  United States        1551000.0
4  United States        181269.0

      Longbusinesssummary     Weight log10_mcap \
0  Apple Inc. designs, manufactures, and markets ...  0.069209  12.585102
1  NVIDIA Corporation provides graphics and compu...  0.059350  12.518356
2  Microsoft Corporation develops and supports so...  0.058401  12.511358
3  Amazon.com, Inc. engages in the retail sale of...  0.042550  12.373837
4  Alphabet Inc. offers various products and plat...  0.042309  12.371368

Outlier_Weight  Outlier_Mcaplog  Outlier_Any
0            True           True        True
1            True           True        True
2            True           True        True
3            True           True        True
4            True           True        True

```

```
[213]: top_marketcap = df.sort_values(by="Marketcap", ascending=False).head(10)
```

```
print(top_marketcap[['Sector', 'Industry', 'Marketcap']])
```


	Sector	Industry	Marketcap
0	Technology	Consumer Electronics	3846819807232
1	Technology	Semiconductors	3298803056640
2	Technology	Software - Infrastructure	3246068596736
3	Consumer Cyclical	Internet Retail	2365033807872
4	Communication Services	Internet Content & Information	2351625142272
5	Communication Services	Internet Content & Information	2351623045120
6	Communication Services	Internet Content & Information	1477457739776
7	Consumer Cyclical	Auto Manufacturers	1351627833344
8	Technology	Semiconductors	1031217348608
9	Financial Services	Insurance - Diversified	978776031232

```
[214]: plt.figure(figsize=(12, 6))
sns.barplot(x='Marketcap', y='Shortname', data=top_marketcap.head(10), palette='coolwarm')
```

```

plt.title("Top 10 Companies in S&P 500 by Market Cap")
plt.xlabel("Market Capitalization")
plt.ylabel("Company")
plt.ticklabel_format(style='plain', axis='x')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



```
[215]: top_price = df.sort_values(by="Weight", ascending=False).head(10)
print(top_price[['Sector', 'Industry', 'Weight']])
```

	Sector	Industry	Weight
0	Technology	Consumer Electronics	0.069209
1	Technology	Semiconductors	0.059350
2	Technology	Software - Infrastructure	0.058401
3	Consumer Cyclical	Internet Retail	0.042550
4	Communication Services	Internet Content & Information	0.042309
5	Communication Services	Internet Content & Information	0.042309
6	Communication Services	Internet Content & Information	0.026581
7	Consumer Cyclical	Auto Manufacturers	0.024317
8	Technology	Semiconductors	0.018553
9	Financial Services	Insurance - Diversified	0.017609

```
[216]: plt.figure(figsize = (10,6))
sns.barplot(x='Weight', y='Industry', data=top_price, palette='bright')
plt.title("Top 10 Companies in S&P 500 by Weight")
plt.xlabel("Weight")
plt.ylabel("Industry")
plt.ticklabel_format(style='plain', axis='x')
```

```
plt.xticks(rotation=45)  
plt.tight_layout()  
plt.show()
```

